

# UCSC Silicon Valley Extension

## Advanced C Programming

### Memory Management

Instructor: Radhika Grover

# Overview

- Heap management functions: malloc, calloc, realloc, free
- How to use the functions to allocate and free memory
- Programming errors when using these functions and how to correct them

# Heap management types

- Types:
  - `void *`: generic pointer to any type and can be converted to other pointer types
  - `size_t`: unsigned integer (can be a typedef of unsigned int, unsigned long, or unsigned long long on different systems)
- Heap management functions: `malloc`, `calloc`, `realloc`, `free`

# Functions malloc and calloc

- `void * malloc(size_t size)`
  - Returns a pointer to the allocated memory of `size` bytes, OR
  - Returns `NULL` if memory could not be allocated.
  - The memory is not initialized.
  - Example, to allocate an array of  $n$  integers:  

```
int *ptr = (int *) malloc(sizeof(int) * n)
```
- `void * calloc(size_t n, size_t size)`
  - Allocates memory for an array of  $n$  elements of `size` bytes each
  - Returns a pointer to the allocated memory.
  - The memory is initialized to zero.
  - Example, to allocate an array of  $n$  integers and set each value to 0:  

```
int *ptr = (int *) calloc(n, sizeof(int))
```

# Initializing memory

- `calloc` initializes memory to 0 but `malloc` leaves this task to programmer
- Can use `memset` in `<string.h>` to initialize memory:
  - `void * memset(void *dest, int c, size_t count)`

*dest* points to the block of memory to be written, *c* is the byte to write into this block, and *count* is the number of bytes to be set. Returns a pointer to *dest*.

# Example: using malloc and memset

```
// HeapMalloc/malloc_example.c
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

int main(void)
{
    int newsize = 5;
    // malloc creates an array to store 5 ints
    int *arr = (int *) malloc(sizeof(int) * newsize);

    if (arr == NULL) {
        // print memory could not be allocated
        printf("Error: could not allocate memory");
        return (-1);
    } else {
        // initialize all bytes in the array to 2
        memset(arr, 2, sizeof(int) * newsize);
        for (int i = 0; i < newsize; i++) {
            printf("%d\n", arr[i]); // print array
        }
        return 0;
    }
}
```

# Function realloc

- `void * realloc(void *ptr, size_t newsize)`
  - Changes the size of the memory block to *newsize* bytes by allocating additional memory if needed.
  - Original contents of memory are not changed.
  - Newly allocated memory (if any) is uninitialized.
  - If *ptr* is not NULL, it must have been returned by an earlier call to `malloc()`, `calloc()`, or `realloc()` - why?
- Questions:
  - What happens if *ptr* is NULL?
  - What happens if *newsize* is 0?

# Function realloc example

```
// HeapRealloc/realloc_example.c
int main(void)
{
    int *arr = NULL;
    int newsize = 5;
    int *ptr = realloc(arr, sizeof(int) * newsize); // realloc creates an array of size 5
    if (ptr == NULL) {
        printf("Error: could not allocate memory");
        return(-1);
    } else {
        arr = ptr; // important because arr points to NULL
        for (int i = 0; i < newsize; i++) {
            arr[i] = i; // put some values in the array
        }
        newsize = 10;
        ptr = realloc(arr, sizeof(int) * newsize); // realloc resizes the allocated array to 10
    }
}
```



# Function realloc example continued

```
    if (ptr == NULL) {  
        printf("Error: could not allocate memory");  
        return(-1);  
    } else {  
        arr = ptr; // important  
        for (int i = 0; i < newsize; i++) {  
            printf("%d\n", arr[i]); // print array  
        }  
        return 0;  
    }  
}  
}
```

# Function free

- `void free(void *ptr)`
  - Frees memory space pointed to by ptr, which must have been returned by a previous call to `malloc()`, `calloc()`, or `realloc()`. (why?)
  - If `free(ptr)` has already been called before, undefined behavior occurs.
  - If ptr is NULL, no operation is performed.

# Exercise: Find the errors

- Find the errors in the code segments:
  1. // program creates a buffer to store 5 ints using malloc and store 10 at  
// each location  
`int *arr = (int *) malloc(sizeof(int) * 5);`  
`arr[0] += arr[0] + 10;`
  2. // program resizes arr to store 5 ints using realloc  
`int arr[1] = {1};`  
`int *ptr = realloc(arr, sizeof(int) * 5);`  
`strcpy(arr, 5);`
  3. // program reads the array size in argv[1] and creates an array of this size:  
`char *arr = (char *) malloc(sizeof(char) * argv[1]);`

# Exercise: Solution

1.

2.

# Exercise: Solution

```
3. // user gives the array size in argv[1]
   int *arr = (int *) malloc(sizeof(int) * argv[1]);
   if (arr == NULL) {
       // print memory could not be allocated
   }
```

There are several problems here:

1. `argv[1]` must never be zero or negative, otherwise the behavior is undefined or may even lead to a buffer overflow
2. The product `sizeof(int) * argv[1]` must not create a value that is greater than the maximum value that is stored in a `size_t` (`SIZE_MAX`); Otherwise, the value will overflow resulting in a smaller number and the buffer that is created will have a smaller size than the size of data that will be stored in it. This can lead of a buffer overflow that may be exploited by a malicious user. For example, if `argv[1]` is `SIZE_MAX`, the result `sizeof(int) * SIZE_MAX` will overflow a `size_t`.

# Exercise: Solution

```
// HeapExercise/solution.c
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    // Run the program with negative and large values of argv[1]; Add check for argc
    long long int len = atoll(argv[1]);
    if (len <= 0) {
        // print error that buffer size must be greater than 0
        printf("%lld Error: Buffer size must be greater than 0", len);
        return (-1);
    }
    else if (len > SIZE_MAX/sizeof(int)) {
        // print error message that the value of len may
        // lead to a buffer overflow
        printf("Error: %lld * sizeof(int) is greater than SIZE_MAX", len);
        return (-1);
    } else {
        // malloc creates an array to store len ints
        int *arr = (int *) malloc(sizeof(int) * len);
    }
}
```

# Exercise: Find the errors

- This program reads a string in argv[1], stores it in a buffer created using malloc and displays it. Find all the errors in the program. (Hint: there are more than five errors.)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>

int main(int argc, char *argv[])
{
    size_t len = strlen(argv[1]);
    char *arr = (char *) malloc(len + 1);
    strncpy(arr, argv[1], len);
    printf(arr);
    free(arr);
    return 0;
}
```

# Exercise: Buffer overflow vulnerability solution

```
// HeapExercise2/solution.c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[])
{
    // Add check for argc
    // calculate the length of the string
    size_t len = strlen(argv[1]); // use strlen_s instead of strlen (why?)
    if (len > SIZE_MAX - 1) {
        // print error message that string size is greater than maximum allowed
        printf("Error: %zu + 1 is greater maximum allowed, buffer overflow may occur", len);
        return (-1);
    } else {
        // malloc creates an array to store a string of length len + 1 for null terminator
        char *arr = (char *) malloc(len + 1);
        if (arr == NULL) {
            // print memory could not be allocated
            printf("Error: could not allocate memory");
            return (-1);
        } else {
            // use a safe function such as strncpy to copy string into arr and then free memory
        }
    }
}
```



# Exercise: Freeing memory incorrectly

- What is the problem in the following program and how will you correct it?

```
// FreeMemory/problem.c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *arr = (char *) malloc(10*sizeof(char));
    if (arr == NULL) {
        // print memory could not be allocated
        printf("Error: could not allocate memory");
        return (-1);
    } else {
        strcpy(arr, "Hello");
        free(arr);
        printf("%s", arr);
        return 0;
    }
}
```

# Exercise: Freeing memory incorrectly solution

- Using strcpy is not a problem in the because the string is known (“Hello”) and fits in the buffer.  
    strcpy(arr, “Hello”);
- The problem is that the contents of memory are being printed out after the free function are called.
- This can lead to a vulnerability where a malicious user can read data that should have been deleted.
- To prevent this, you should set the pointer to NULL after free is called.  
    free(arr);  
    arr = NULL;

# Checking for memory leaks

- Use a tool to determine if heap memory (this is memory allocated with malloc, calloc, realloc) has been leaked .
- Valgrind is available for memory profiling and runs on several Linux systems. Visual Studio on Windows machines has a built in memory profiler.
- Download valgrind from <http://valgrind.org> and install it. You can run it on the command line for the executable Stack as:
- `valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all ./Stack`

# Further reading

- Reference: <https://en.cppreference.com/w/c/memory>
- Robert Seacord. Secure Coding in C and C++.